AI Agent Skills and MCP: Building the Right AI Agent Pattern

Stephen Simon

Chapter 3: Building AI Workflows with MCP

What MCP Solves That Skills Cannot

Compare the Skill-only and MCP reports to see exactly which problems access solves

You now have two reports for the same launch. One was produced by the `launch-risk-review` Skill from pasted documents. The other was produced through MCP from Orbit Labs' live systems, with no Skill. In this video you compare them, trace a finding back to its source, and practice deciding when a workflow needs MCP rather than a better Skill.

Before you start

You need two files in `orbit-launch/reports/`:

- `skill-only.md` from video 02_04.
- `mcp-only.md` from video 03_02.

The `orbit` server should still be connected, with the new ticket from video 03_02 still in the data. If you're missing a report, go back and repeat the step that saves it. If you ran `npm run reset` since video 03_02, run `npm run demo:new-ticket` in `orbit-backend` to add the ticket again.

Step 1: Open both reports side by side

Open `reports/skill-only.md` and `reports/mcp-only.md` in your editor, side by side. Both answer the same question: what could put the Orbit 2.0 launch at risk?

Step 2: Compare what each report knows

Go topic by topic and compare what each report says:

- **Duplicate charges.** Skill-only: "four customers affected," with ticket priorities listed as missing information. MCP: four **P1** tickets, each with an ID, open between 8 and 13 days.
- **PayGate.** Skill-only: found through the vendor email, with a conflict noted against the status update. MCP: found in the notes on task T-104, which is marked as blocked.
- **Pen test.** Skill-only: no owner as of the October 3 meeting. MCP: task T-117 still has no owner today.
- **SSO outage.** Skill-only: not mentioned anywhere. It happened after every document was written. MCP: ticket TCK-2031, six customers, P1.

- **Risk register.** Skill-only: can't be updated. MCP: the `create_risk_item` tool is available, even though you haven't used it yet.

Notice that the Skill-only report wasn't wrong about everything. It caught the PayGate slip and the pen test, because someone had written them down. Documents work when they're recent and accurate. The gaps appear where information was summarized by hand, changed after the last document, or lives in a system nobody exported.

Step 3: Trace a finding back to its source

Pick one finding and check it against the system it came from. Send:

```
Using the orbit server, show me the details for ticket TCK-1987.
```

The agent calls `get_ticket` and returns the ticket: a P1 opened on September 22 by Brightwave Studio, open for 13 days. The evidence in the MCP report points to something you can verify. The Skill-only report could only point to a sentence in a meeting note.

Step 4: Decide which requests need MCP

For each request below, decide: does it need MCP, a Skill, or neither? Then check your answer.

Request 1: "What did the team agree in the October 3 go/no-go meeting?"

No MCP needed. The meeting notes are the source, and they don't change. A document in the workspace is enough.

Request 2: "Is the pen test booked yet?"

MCP. The answer lives in the project tracker and may have changed since the last document.

Request 3: "How should we rate a cluster of P2 tickets that have been open for ten days?"

A Skill, not MCP. This is about how the team judges risk. No external data would answer it.

When MCP is the better solution

MCP is the better fit when the gap in a workflow is one of these:

- **Freshness.** The answer depends on what's true now, and the information changes often.
- **Location.** The information lives in another system, and someone would otherwise export or copy it.
- **Fidelity.** Details get lost when people summarize by hand, like four P1 tickets becoming "a couple of reports."
- **Traceability.** You need evidence that points back to a source of truth.
- **Action.** The workflow needs to create or update something in another system.

If the gap is how the work should be done (the rules, the format, the judgment), MCP doesn't close it. Look at the MCP report again: it has the better facts, but its format and severity labels were the agent's choice. That's the subject of the next video.