AI Agent Skills and MCP: Building the Right AI Agent Pattern

Stephen Simon

Chapter 4: Combining Agent Skills and MCP

Skills, MCP, or Both?

Decide which pattern a workflow needs, and spot the signs you picked the wrong one

You've now built the Orbit Labs launch review four ways: as a prompt, as a Skill, with MCP, and with both. In this video you step back from the project and apply what you learned to other workflows, using the decision framework from Chapter 1.

The two questions

Start with the workflow, not the technology. Ask:

1. Does the workflow need reusable know-how about how to do this type of task?
2. Does it need information or actions from outside the agent's current context?

The answers point to one of four patterns:

- **No to both: a prompt.** A one-time request where the agent already has what it needs.
- **Yes to the first only: a Skill.** Reusable know-how, applied to information the agent already has.
- **Yes to the second only: MCP.** Access to external data or actions, with no special process to follow.
- **Yes to both: Skills and MCP together.** Reusable know-how applied to live systems.

The framework applied to Orbit Labs

The launch review answers yes to both questions:

- **Reusable know-how:** the severity rubric, the filters, the evidence rules, the report template, and the rules for writing to the register.
- **External information and actions:** the milestones, tasks, and tickets in Orbit Labs' systems, and the shared risk register.

That's why each partial version fell short. The Skill alone had stale data and no way to act. MCP alone had fresh data but inconsistent judgment and no rules for writing.

Practice: decide for each workflow

These four scenarios are the ones discussed in the video. For each one, answer the two questions, decide on a pattern, then check your answer.

Scenario 1: Review a pull request against the team's checklist: security, breaking API changes, test coverage, and blocking versus optional comments. The changes are already in the local repository.

A Skill. The checklist and comment format are reusable know-how, and the code is already in the workspace. If the agent also had to fetch the pull request from GitHub or post comments there, it would become Skills and MCP together.

Scenario 2: Tell me which of my meetings tomorrow conflict.

MCP. The calendar is external and changes constantly. There's no special process to encode.

Scenario 3: Translate this paragraph into Spanish.

A prompt. It's a one-time request, there's nothing to reuse, and nothing to fetch.

Scenario 4: During an incident, classify severity using our criteria, pull the recent error logs, follow the escalation runbook, and open an incident ticket.

Skills and MCP together. The criteria, runbook, and report format are know-how. The logs and the ticketing system are external.

Try two more

These aren't covered in the video.

Scenario 5: Write commit messages that follow our Conventional Commits rules.

A Skill. Reusable know-how, and the changes are already in the local repository.

Scenario 6: Find the open bugs tagged "billing" and summarize them.

MCP. Unless the team has a specific way bug summaries should be written. Then it becomes Skills and MCP together.

Signs you picked the wrong pattern

Workflows change, and the right pattern can change with them. Watch for these symptoms:

- **People paste exports from another system into the chat every time.** The workflow needs MCP.
- **The same rules are copied into many prompts and slowly drift apart.** The workflow needs a Skill.
- **Results through MCP vary between people and between runs.** Add a Skill with the process.
- **Team rules are written into MCP tool descriptions.** Move them into a Skill.
- **A Skill contains a copy of data that goes stale.** Fetch that data through MCP instead.
- **A one-off question became a Skill nobody reuses.** A prompt was enough.

Keep in mind

- **Both isn't the default.** Add a Skill or MCP when the workflow needs it, not because it's available. Many tasks are best served by a prompt.
- **Workflows grow.** It's normal to start with a prompt, extract a Skill once instructions start repeating, and add MCP once the agent needs access. That's the path the Orbit Labs review followed.
- **Look at the parts.** A single workflow can mix patterns. In the launch review, judging risk is know-how, while fetching tickets and writing to the register is access.
- **Weigh the risk of actions.** If MCP gives the agent write access to a shared system, the rules for when to write are know-how. That usually means a Skill as well.